



# SeLeP: Learning Based Semantic Prefetching for Exploratory Database Workloads

Farzaneh Zirak

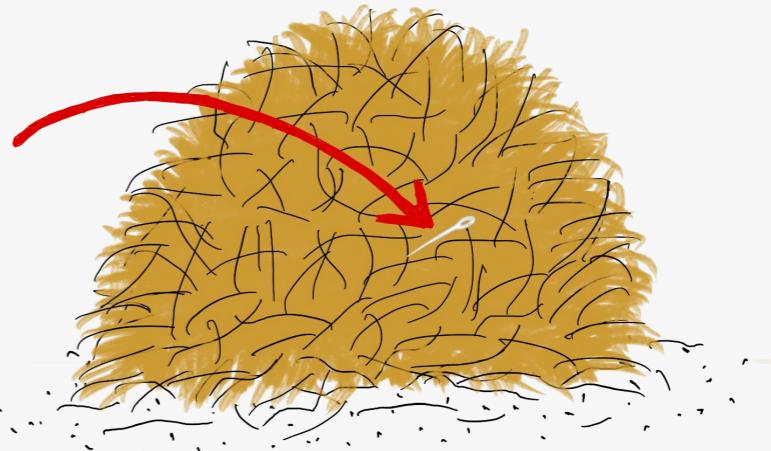
Supervisors:

Renata Borovica-Gajic, Farhana Choudhury

# data exploration

- Searching for insights buried in vast amount of data
- Data being collected at a striking rate

Rising need for effective data exploration



- Challenges and requirements:

**No a priori knowledge**

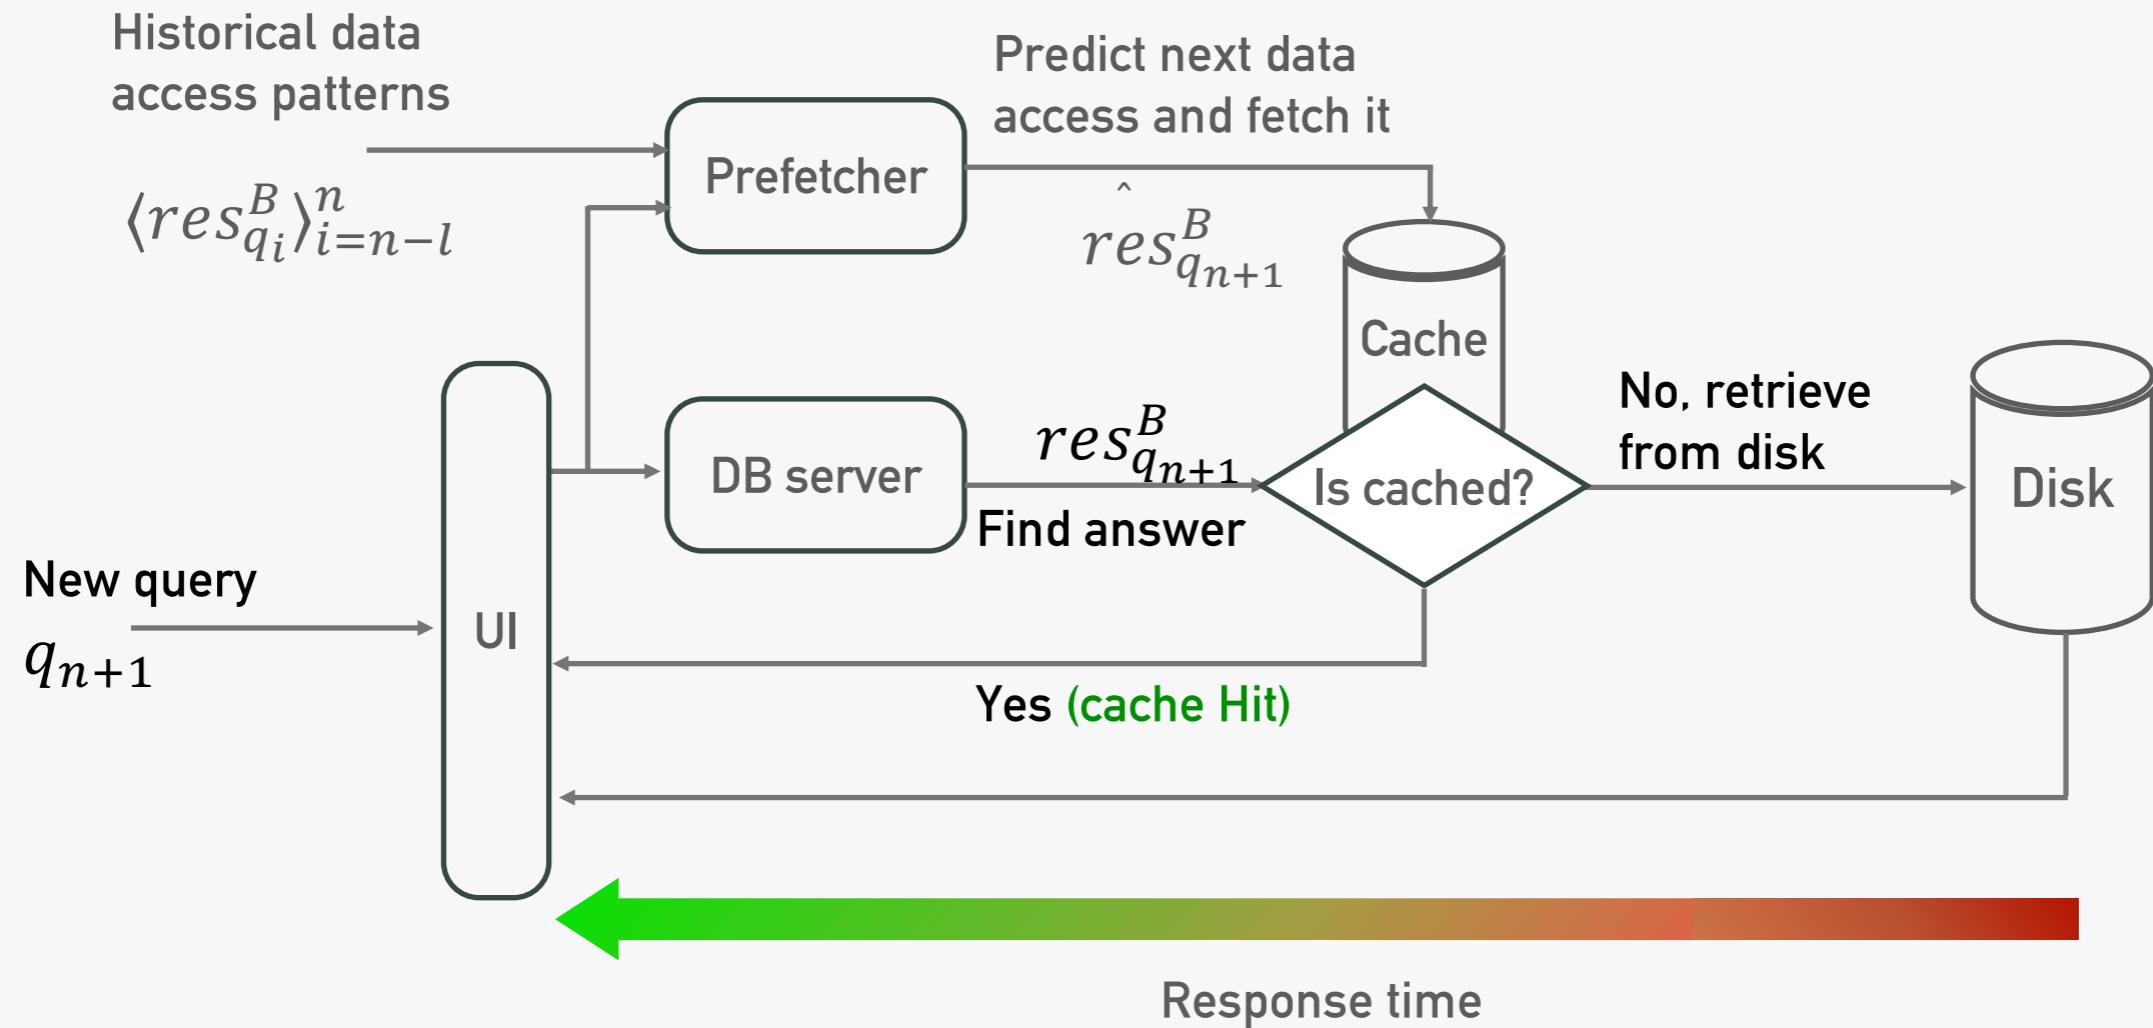
**Interactive**

**Non-expert user**

**Adaptive**

# Insight

- Increase interactivity by reducing response time using **prefetching** technique
- Predict the upcoming workload and fill the cache efficiently with data to be accessed



# Objectives

Semantic based

Make prefetch decision considering actual data values (instead of using logical Block Addresses)

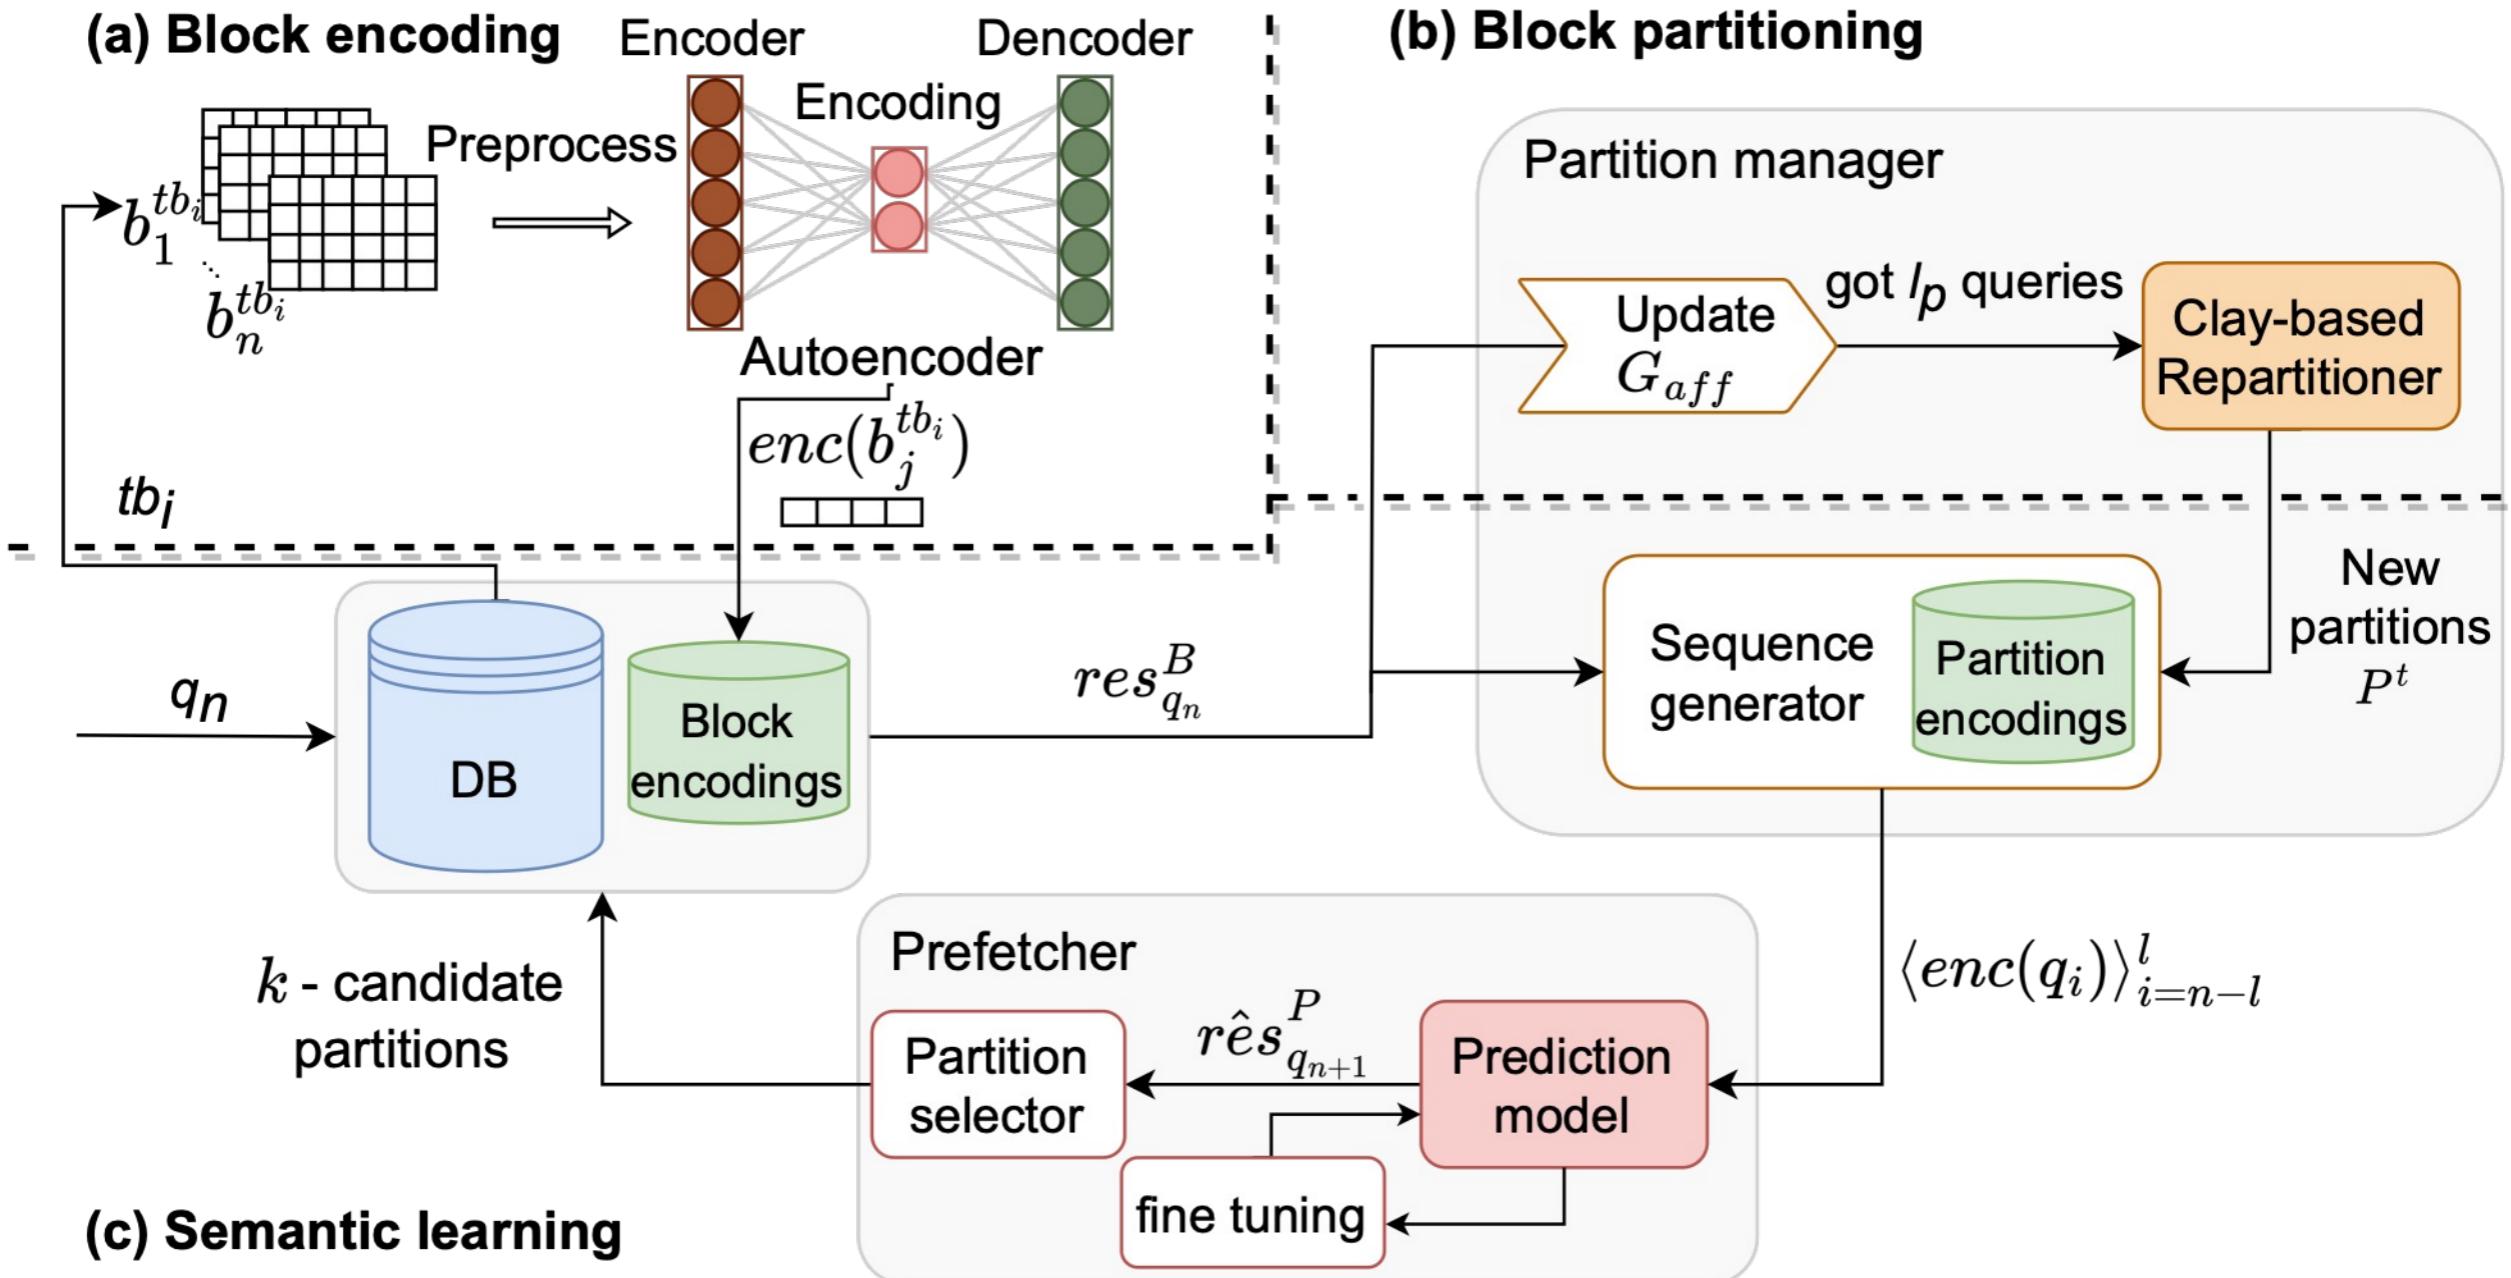
Suitable for various workloads

Support both navigational and SQL-based workloads

Adaptive

Can adapt to workload shifts

# SeLeP Overview [VLDB'24]



# Block encoding

Semantic based

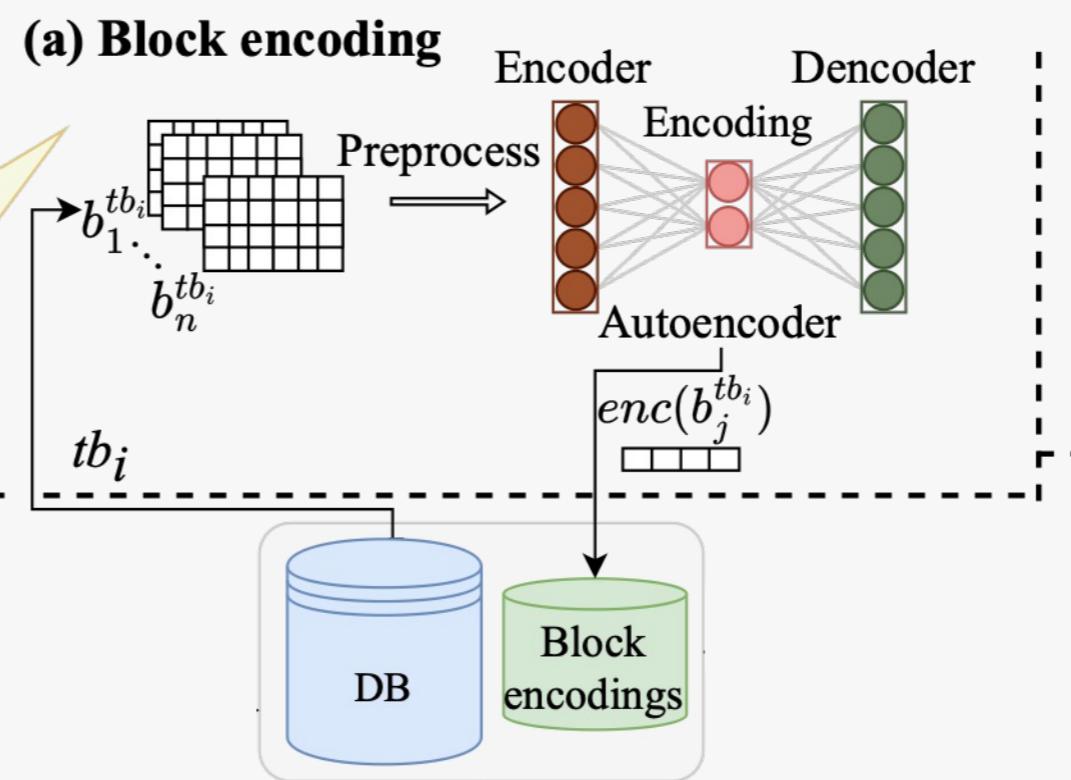
- Each block contains hundreds of values
- Cannot detect the connection between data values in a sequence of blocks
- Need a concise representation for each block which captures the essential and distinctive characteristics of its data

Consider each block of data as a big matrix and encode them into a vector of length  $l_b$  using Autoencoders

**(eg-1)**  $b_3^{tb_2}$ ,  $tb_2 = ObjAll$

g	r	type	size
30.9	4.6	'star'	2.85
31.4	4.5	'star'	1.86
30.7	24.2	'galaxy'	92.62
32.9	4.3	'star'	2.32

avg(r) near min  
avg(g) near max  
Mostly stars  
There are Galaxies  
small sizes



# Block encoding

Semantic based

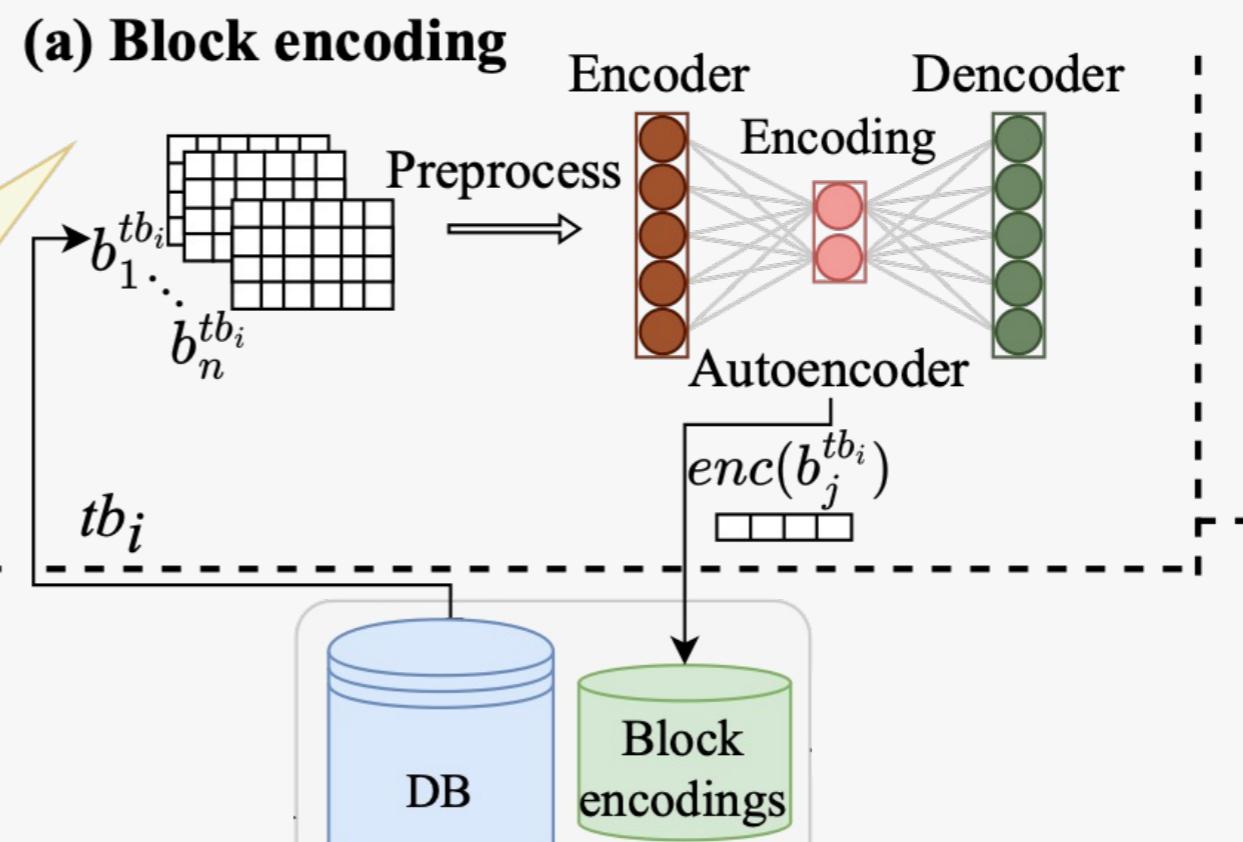
- Encoder and decoder have same layer structure consisting of two dense layers
- Preprocessing includes dimensionality reduction using PCA, string to number conversion using word2vec, and data normalisation

**(eg-1)**  $b_3^{tb_2}$ ,  $tb_2 = ObjAll$

g	r	type	size
30.9	4.6	'star'	2.85
31.4	4.5	'star'	1.86
30.7	24.2	'galaxy'	92.62
32.9	4.3	'star'	2.32

avg(r) near min Encodes small sizes

avg(g) near max Mostly stars There are Galaxies



# Query encoding

Support various workloads

- Visual queries are easy, a few adjacent blocks get accessed
- SQL queries are harder to predict, a set of blocks are getting accessed with various size from different part of the data space
- Need a better representation for queries:

Represent queries by aggregated semantics of its accessed data

- Instead of directly using the blocks  $\langle res_{q_i}^B \rangle_{i=n-l}^n$ , we can input a sequence of encodings
- Using query encodings, the semantic prefetching problem can be transferred to a time series forecasting problem using a multi-label classification model:

query encoding sequence  $\langle enc_{q_i} \rangle_{i=n-l}^n$  as input and  $res_{q_{n+1}}^B$  as the output

# Query encoding

Support various workloads

- Visual queries are easy, a few adjacent blocks get accessed
- SQL queries are harder to predict, a set of blocks are getting accessed with various size from different part of the data space
- Need a better representation for queries:

Represent queries by aggregated semantics of its accessed data

- Instead of directly using the blocks  $\langle res_{q_i}^B \rangle_{i=n-l}^n$ , we can input a sequence of encodings
- Using query encodings, the semantic prefetching problem can be transferred to a time series forecasting problem using a multi-label classification model:

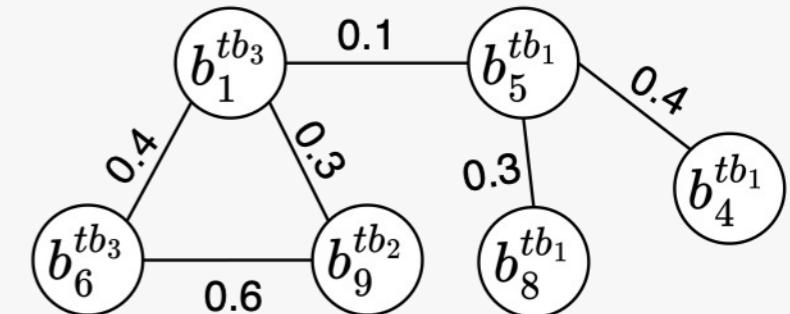
query encoding sequence  $\langle enc_{q_i} \rangle_{i=n-l}^n$  as input and  $res_{q_{n+1}}^B$  as the output

Substantial number of labels in  
the classification model

# Block partitioning

- To reduce number of classes, we cluster blocks frequently accessed together into partitions, which are more sizeable sets
- Construct affinity graph  $G_{aff}$ , with nodes representing the blocks and edge weights illustrating co-accessing rate
- Cluster nodes using graph partitioning and load balancing methods
- Employed Clay partitioning which tries to minimize the inter cluster edges

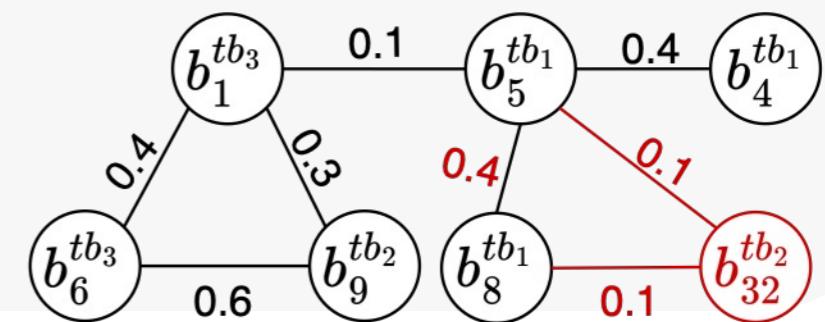
Previous  $G_{aff}$  :



$res_{q_n}^B$  :

$b_5^{tb_1}$     $b_{32}^{tb_2}$     $b_8^{tb_2}$

Updated  $G_{aff}$  with  $l_p = 10$  :

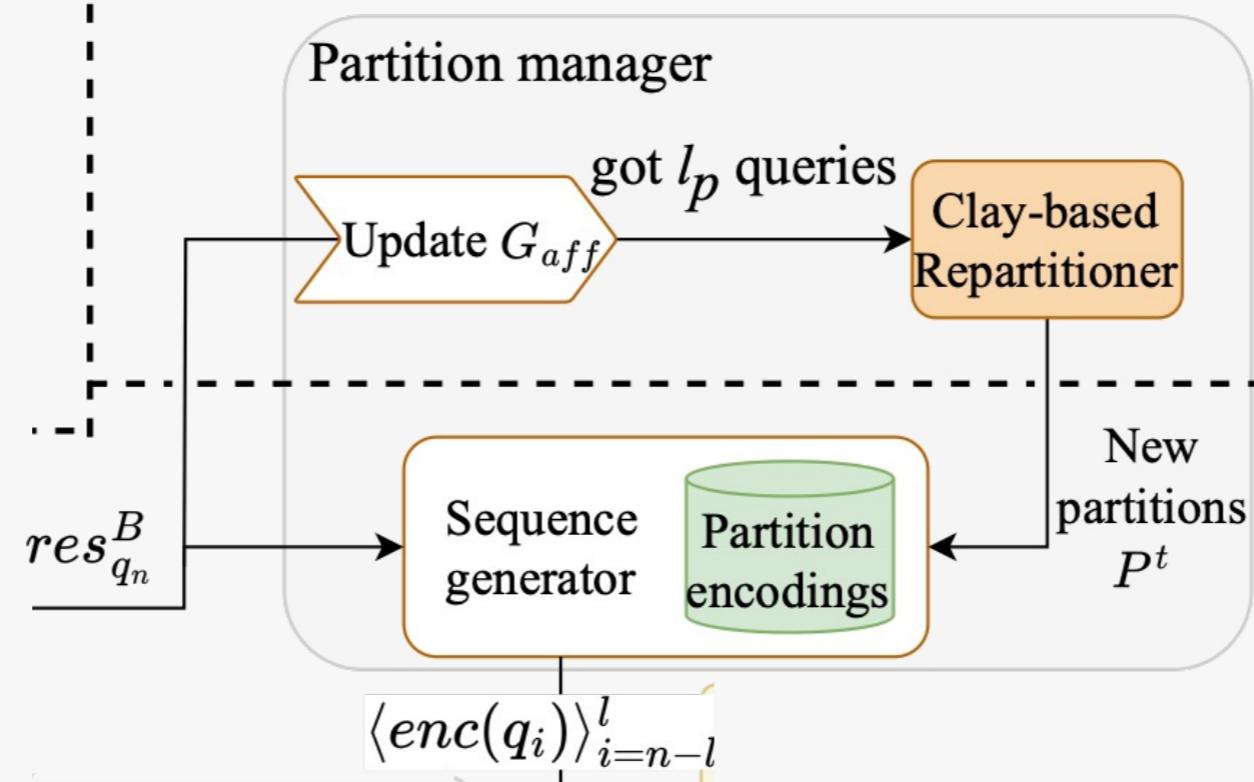


# Block partitioning

- Aggregate block encodings to calculate partition encodings
- Instead of block address, use the partition number to refer to a data access
- Now we can input sequence of query encoding  $\langle enc_{q_i} \rangle_{i=n-l}^n$  calculated by aggregating the partition encodings and get  $res_{q_{n+1}}^P$  as the output

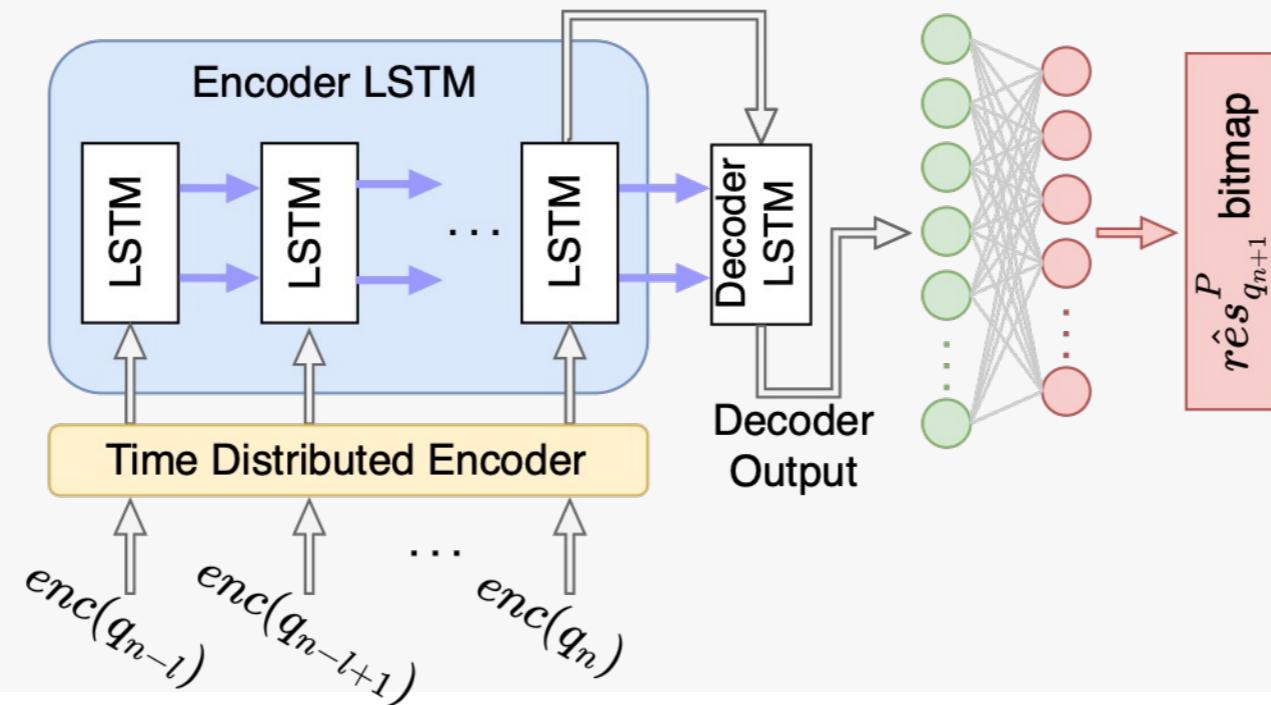
$tb_1$	0	0	0	0	0
$tb_2$	.89	.21	.79	.68	.24
$tb_3$	0	0	0	0	0
$tb_4$	.4	.13	.33	.19	.81

**(b) Block partitioning**



# Semantic learning

- Learn data access pattern and fine tune the model with new workloads Adaptive
- The sequence of query encoding matrices are compressed via a time-distributed layer with 128 units
- The compressed vectors are fed into the encoder and decoder, both designed with a single-layer LSTM with 32 units
- The decoder's output further passes through two dense layers, each containing  $|P|$  units
- The final dense layer generates a bitmap vector with a length of  $|P|$ , representing the probabilities of each partition being accessed in the subsequent query.

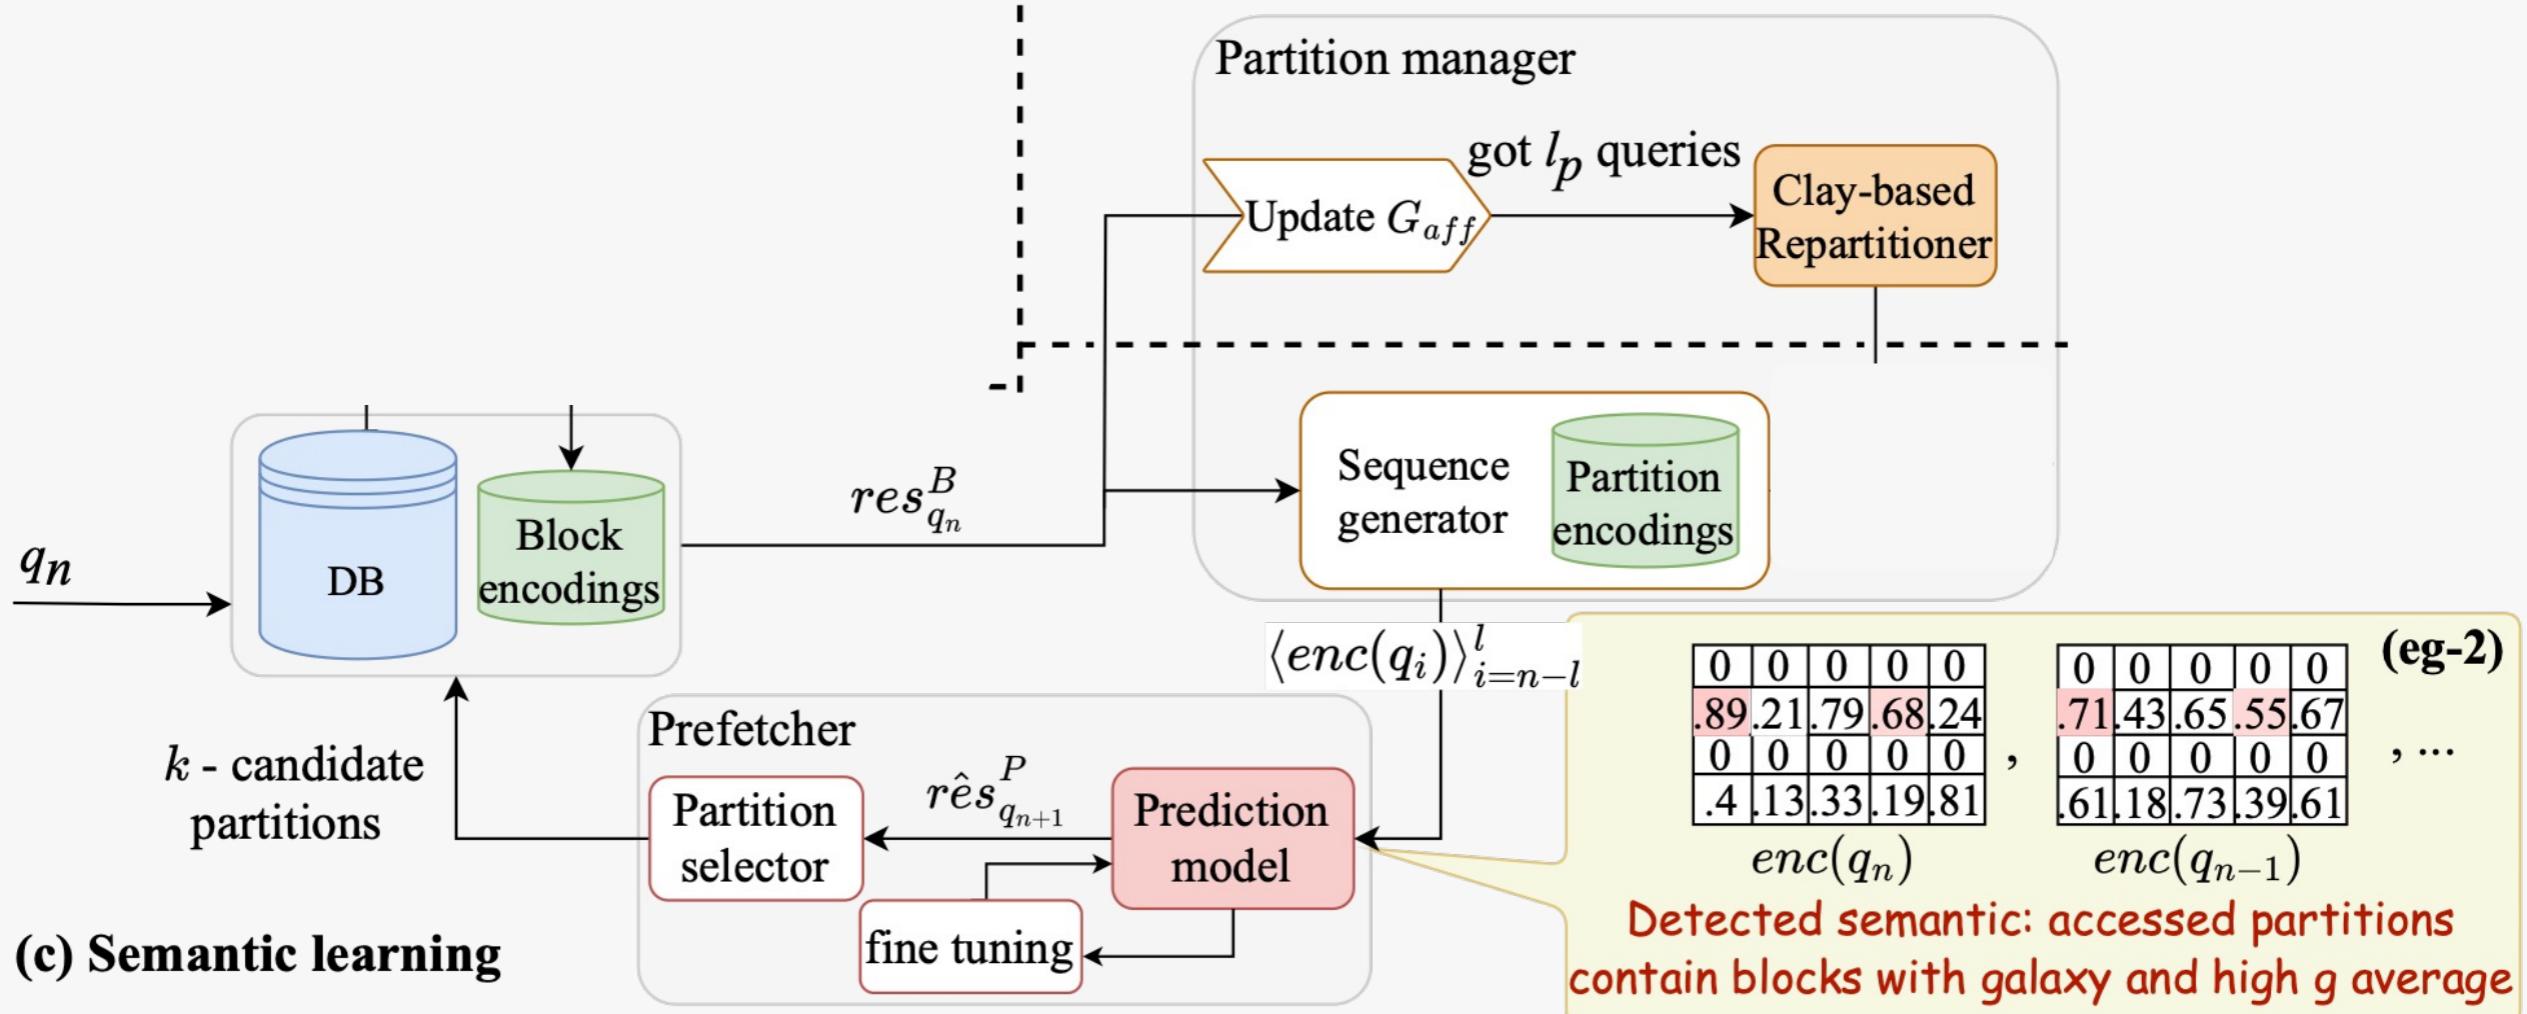


# Semantic learning

Adaptive

Continues repartitioning and fine tuning using the most recently received workload.

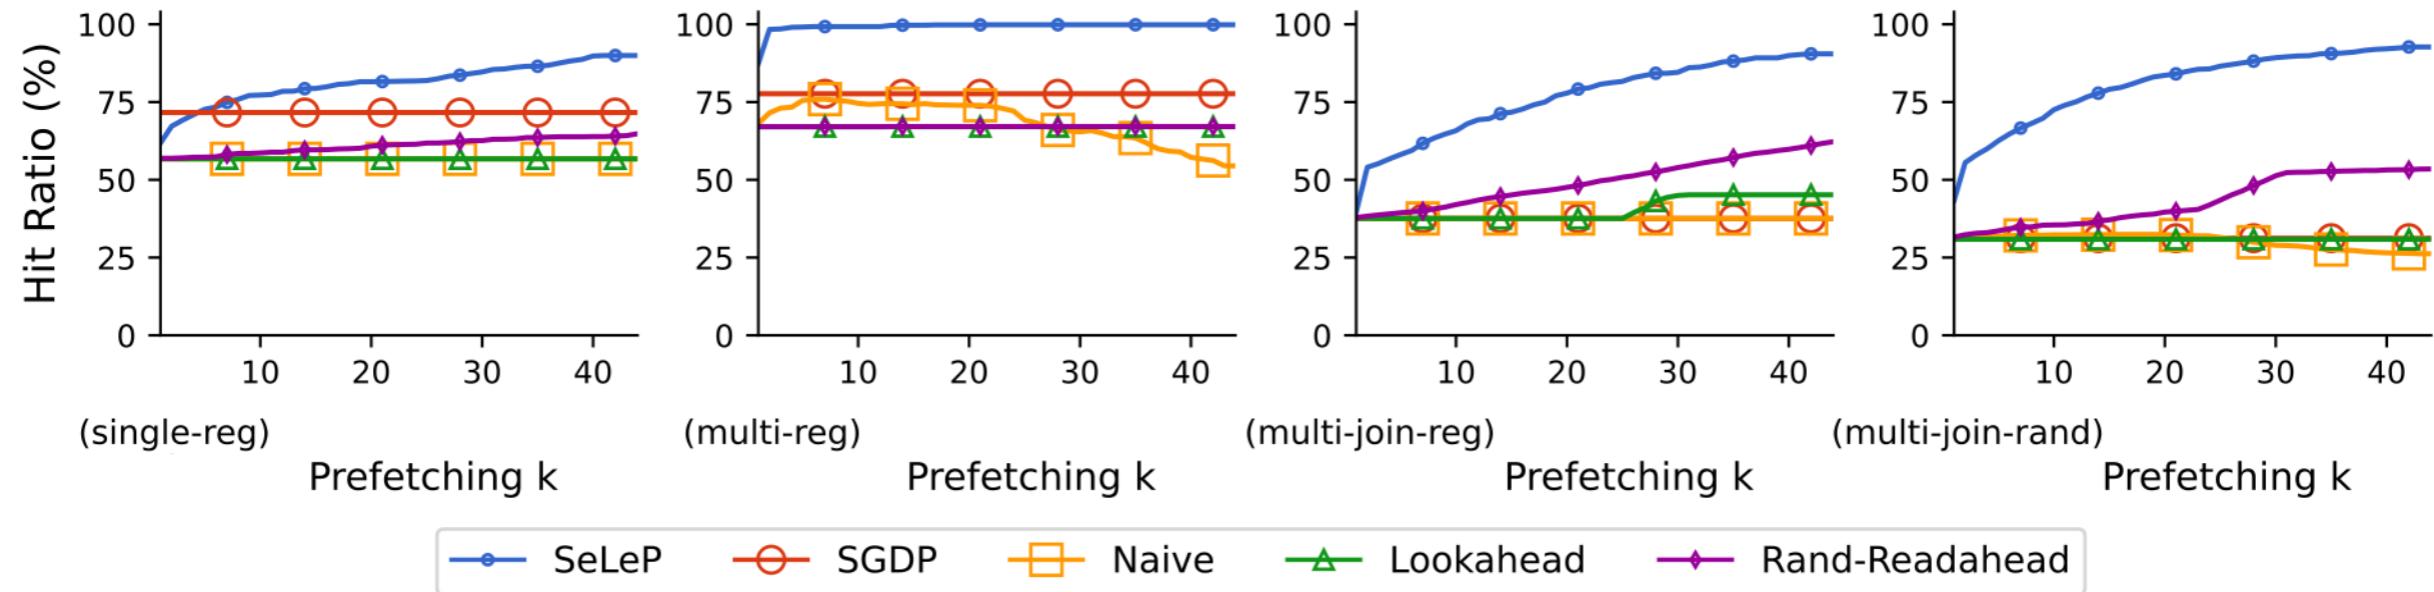
**(b) Block partitioning**



# SeLeP in Action - Hit Ratio

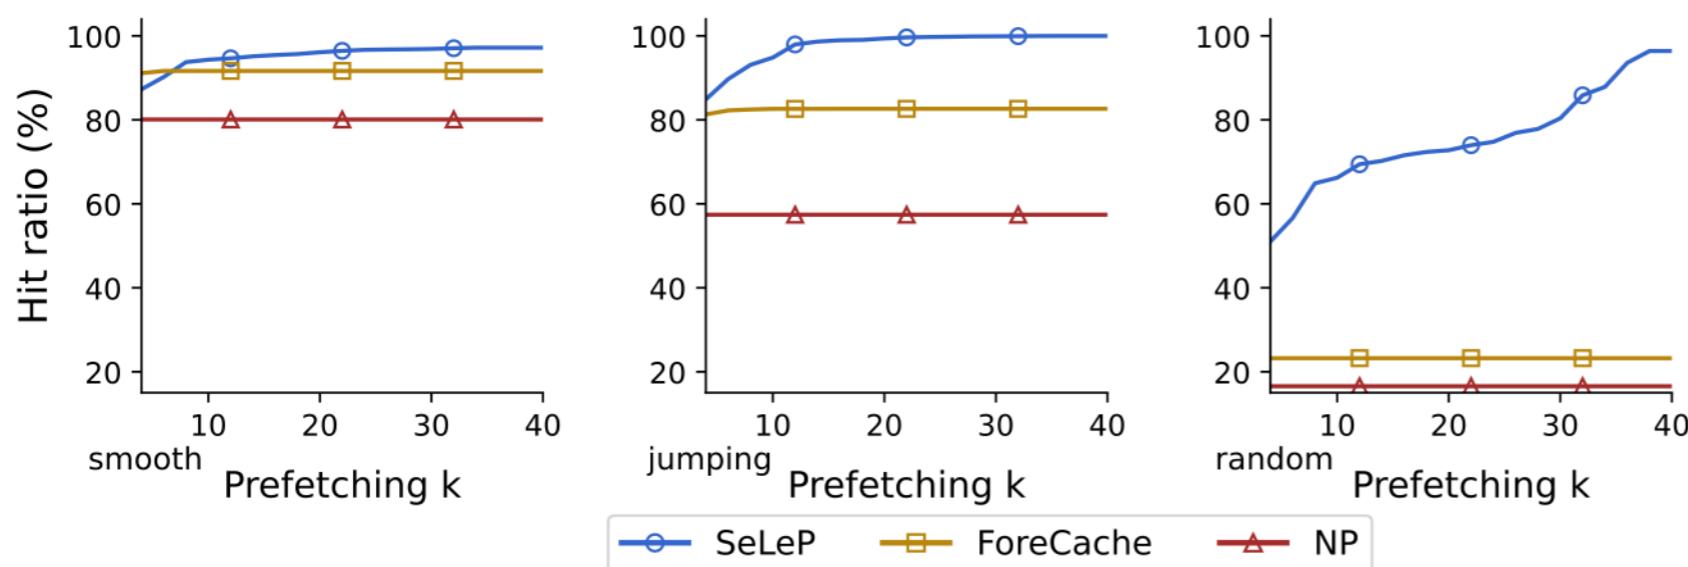
**Setting:** 16GB SDSS DR7, prefetch size =  $k \times 128$  block, 4GB cache

**Queries:** Predictable(reg) and unpredictable(rand) single-table, multi-table, and containing join SQL workloads



**Setting:** 8GB SDSS DR7, prefetch size =  $k \times 64$  block, 2GB cache

**Queries:** Smooth walk, jumping and random navigational workloads



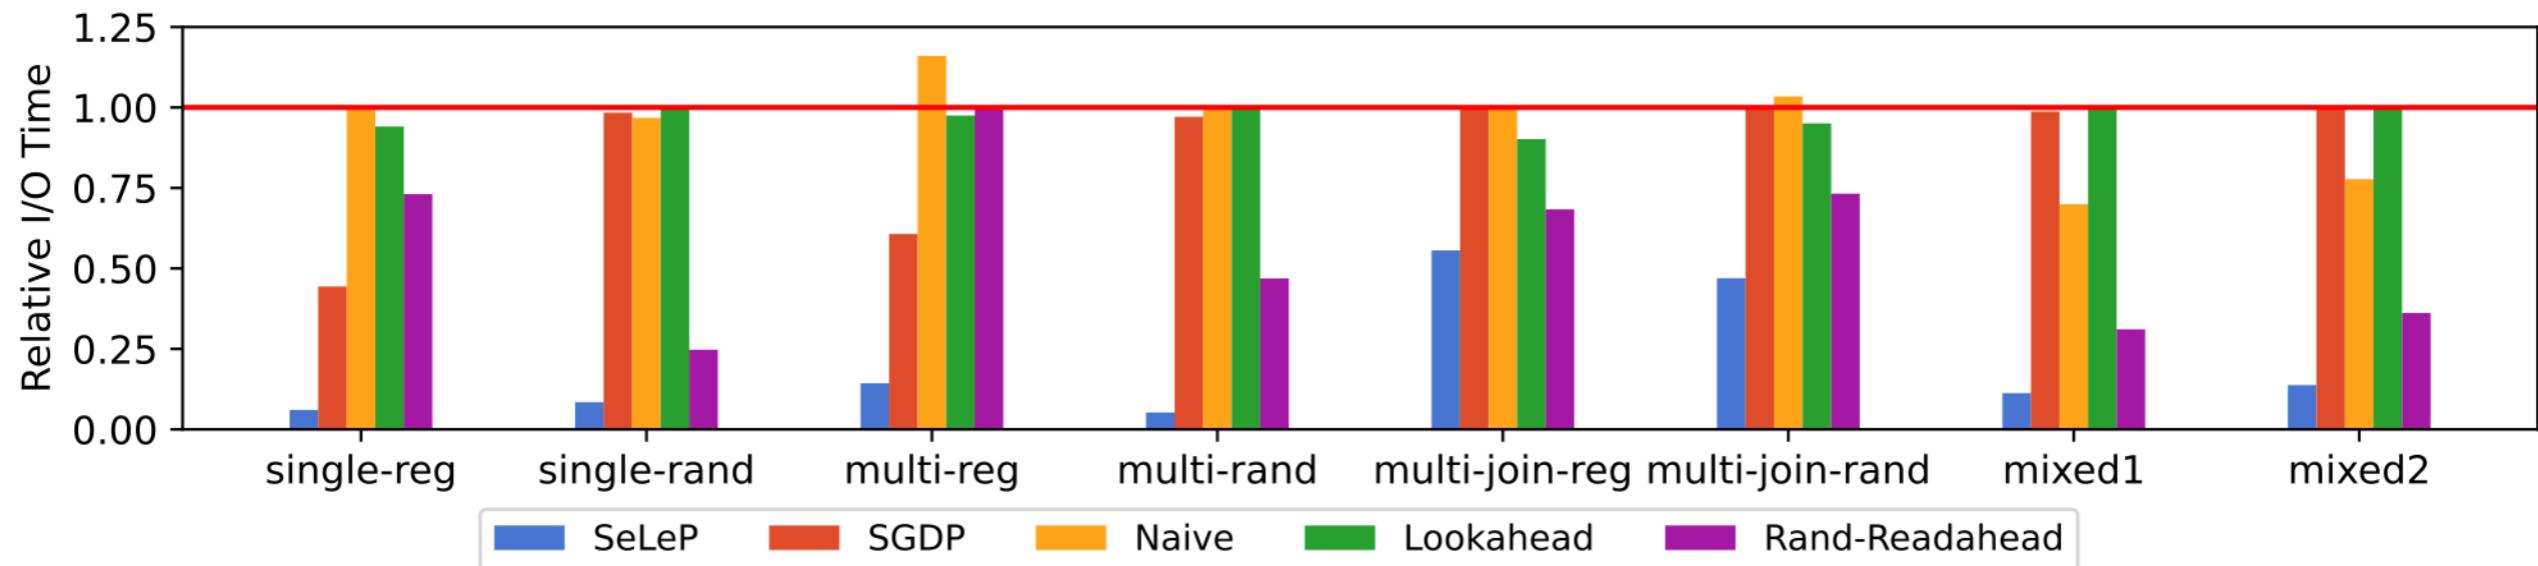
$$\text{Hit ratio} = \frac{\#hits}{\#hits + \#misses}$$

- **95% hit ratio on average**
- **Up to 40% improvement** compared to the traditional and state-of-the-art prefetchers

# SeLeP in Action - I/O Time

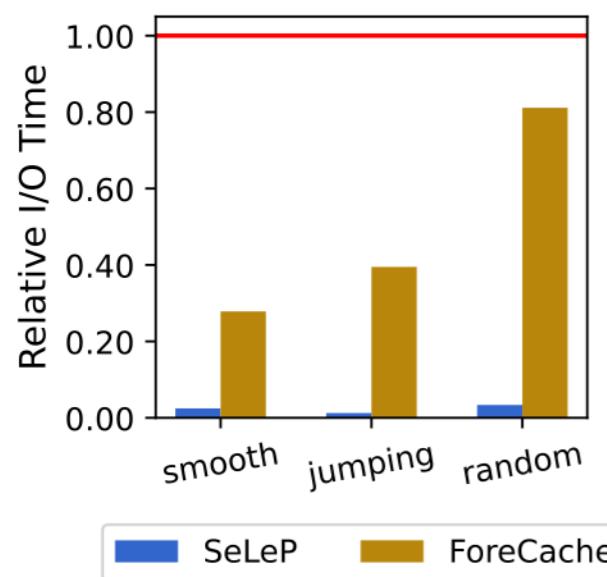
**Setting:** 16GB SDSS DR7, prefetch size =  $k \times 128$  block, 4GB cache

**Queries:** Predictable(reg) and unpredictable(rand) single-table, multi-table, and containing join SQL workloads



**Setting:** 8GB SDSS DR7, prefetch size =  $k \times 64$  block, 2GB cache

**Queries:** Smooth walk, jumping and random navigational workloads

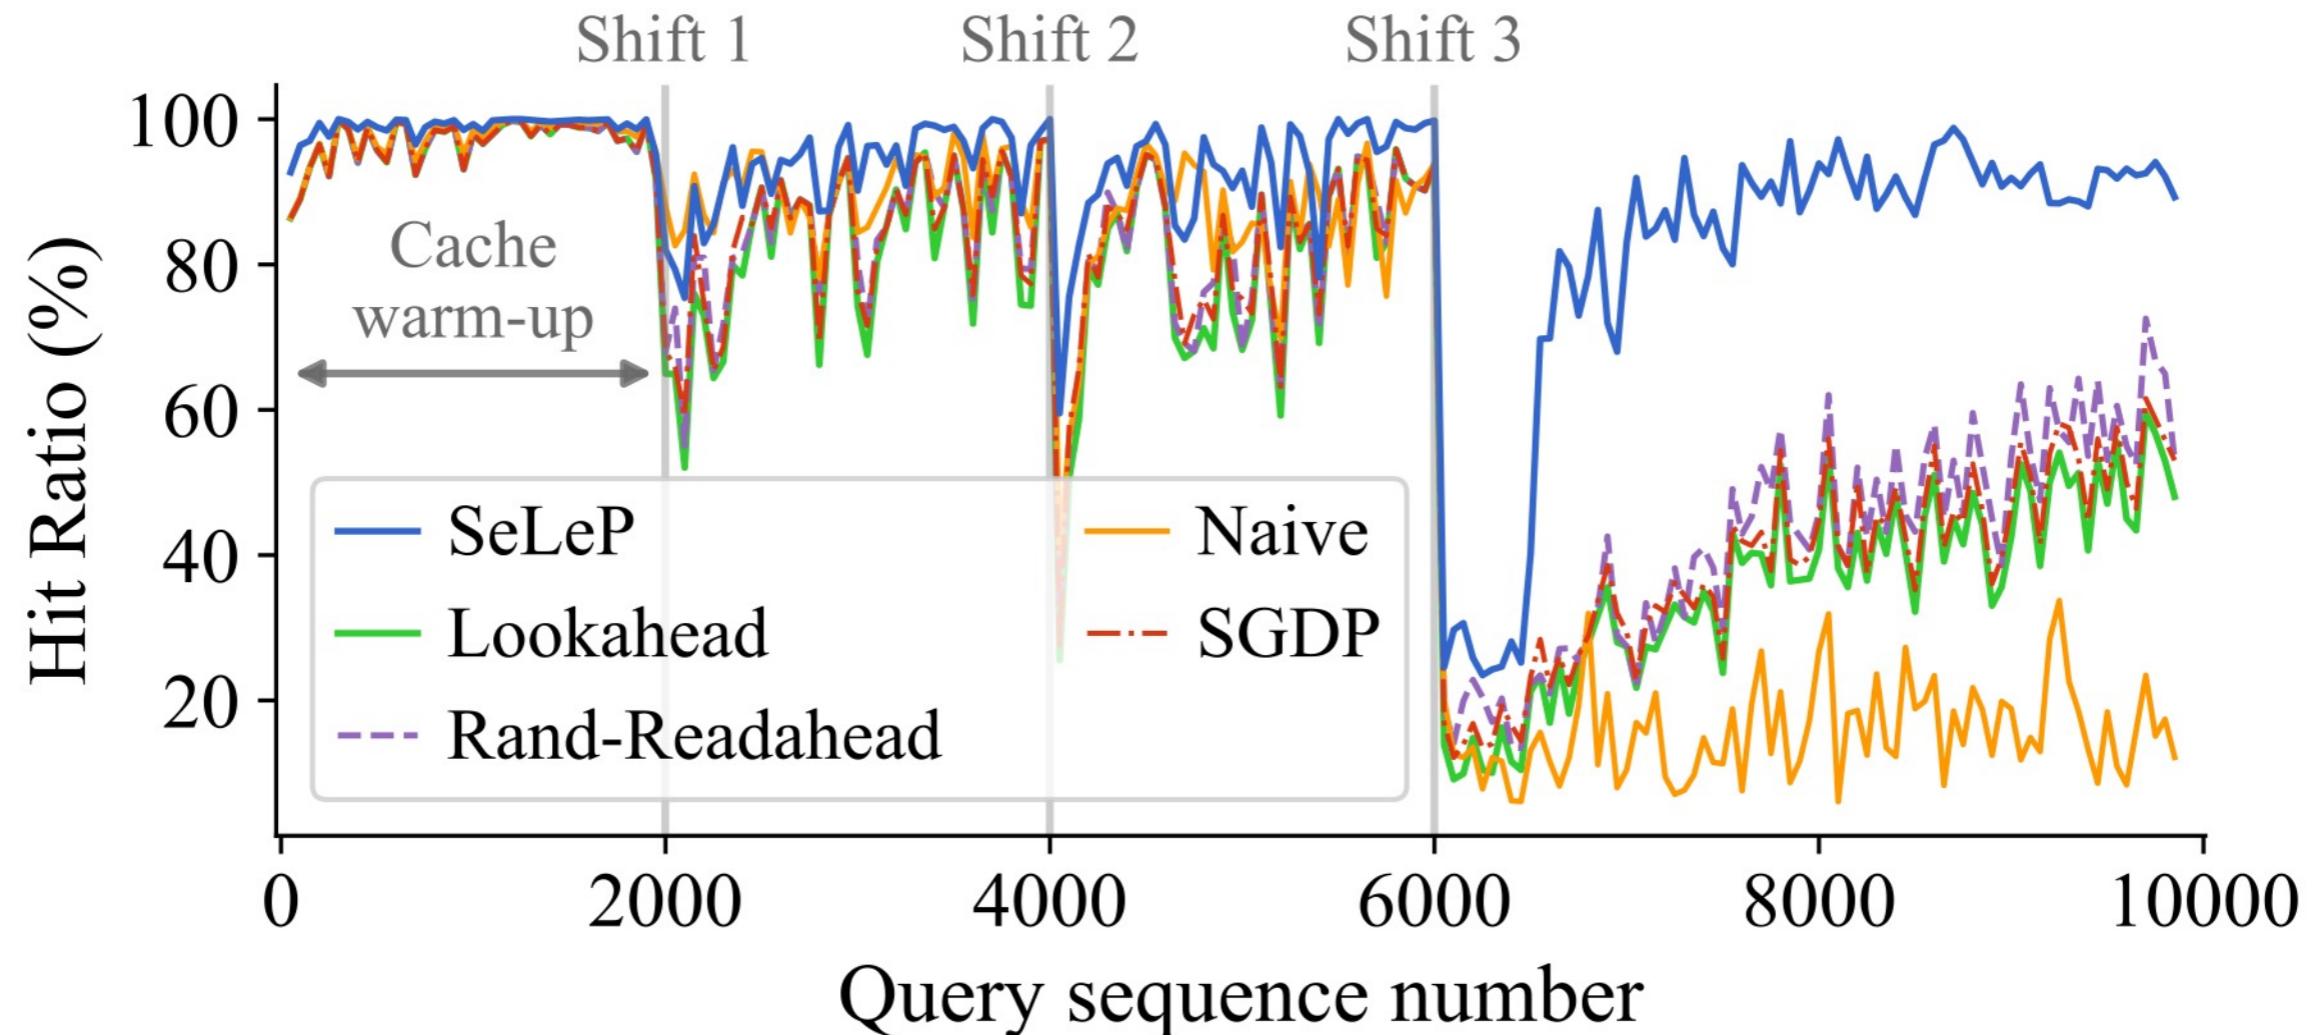


- **80% I/O time reduction on average**
- Up to **45% improvement** compared to the traditional and state-of-the-art prefetchers

# SeLeP in Action - Adaptivity

**Setting:** 16GB SDSS DR7, prefetch size =  $k \times 128$  block, 4GB cache

**Queries:** Shifts at sequence num = {2000, 4000, 6000} with novel query template and new accessed data



# Summary of SeLeP

- Encodes blocks into vectors and **extract their semantics** using AutoEncoders
- **Dynamically partitions blocks** frequently accessed together
- **Learns partition access pattern** using LSTM model

**Improves hit ratio** up to 40% and **reduces I/O time** up to 45% compared to traditional and state-of-the-art prefetchers

**Enables high performance prefetching for SQL workloads**

**Publication:** The paper proposing SeLeP is accepted to be published in VLDB 2024 conference.



THE UNIVERSITY OF  
MELBOURNE

# Questions?

# Thank You